

# In-memory Spatial-Aware Framework for Processing Proximity-Alike Queries in Big Spatial Data



Isam Mashhour Al Jawarneh, Paolo  
Bellavista, Antonio Corradi, Luca  
Foschini, Rebecca Montanari,  
Andrea Zanotti

DISI, University of Bologna, ITALY  
{isam.aljawarneh3, paolo.bellavista,  
antonio.corradi, luca.foschini,  
rebecca.montanari}@unibo.it,  
andrea.zanotti@studio.unibo.it}

Monday, 17<sup>th</sup> September 2018, IEEE CAMAD'18

# Quality of Service Aware Data Stream Processing for Highly Dynamic and Scalable Applications




Isam Mashhour Hasan Al Jawarneh  
PhD Candidate, Cycle XXXII

DISI, University of Bologna, ITALY  
[isam.aljawarneh3@unibo.it](mailto:isam.aljawarneh3@unibo.it),

Supervisor:  
**Prof. Rebecca Montanari**

PhD program coordinator:  
**Prof. Davide Sangiorgi**


***PhD. Esame finale anno 2020, 2nd***



# Agenda

---

- **Smart City and Big Data Context**
  - Geolocated big data
  - MapReduce
  - Spark and GeoSpak
- **Spatial-Aware Big Data Management Strategies**
  - Self-Adaptable Spatial-Aware Partitioner (SASAP)
  - Spatial-aware query optimizations
- **Experimental Results**
  - Partitioning results
  - Proximity, containment and density based clustering
  - query performance results
- **Conclusions & Ongoing Works**



# Smart City and Big Data Context

## Smart City

Advanced technological services



## Geographic Big Data

Huge amount of information


## Mobile Sensing


Automatic collection of data



# Geo-located Big Data


- Large amounts of geolocated data, **exceeding processing capability** of traditional database management systems
- Characteristics
  - **Volume:** amount of data
  - **Velocity:** streaming data
  - **Variety:** multiple sources, heterogeneous data
- Other characteristics
  - **Veracity:** uncertainty degree
  - **Variability:** possible inconsistency
  - **Complexity:** difficult to establish connection and relationships between data






# Advanced Geospatial Queries & Clustering DBSCAN Algorithm

- **Eps ( $\epsilon$ )** = max distance
- **MinPts** = minimum number of points inside radius  $\epsilon$



- **Core points (cp):**  $\#\{X: \text{dist}(X, Y) \leq \epsilon\} \geq \text{MinPts}$
- **Cluster:**  $\{X: \text{dist}(X, cp) \leq \epsilon\} \cup \{cp\}$





# MapReduce

Programming paradigm  
for computing and  
aggregating **large**  
**amounts of data**


Functions:


- **Map**: for each input data, returns a **key-value pair**
- Intermediate **grouping** and **sorting** by key step
- **Reduce**: final aggregation step





# DBSCAN-MR







# Apache Spark


Data stored in memory: analyses on in-memory data and use of **Resilient Distributed Dataset (RDD)**

Hadoop Distributed  
File System




Create  
RDD






# GeoSpark

- **Geospatial-Aware solutions**
  - **GeoSpark** : runs exactly as Spark, but with the awareness for geospatial data, consisting of three layers; Spark, spatial RDD and spatial query processing layers
  - Some other competitors are based on Hadoop like SpatialHadoop and Hadoop-GIS
- **Problem**
  - ***Does not include an integrated support for clustering methods***, or customizable modules for specific application requirements
  - Do not consider specific requirements like ***domain-specific data load balancing for optimizing clustering algorithm execution***




**GeoSpark**



# Main Goals and Requirements

- Realize a **transparent partitioning** support of data, based upon **location information**
- Extend transparently the GeoSpark support
- Realize an effective support to proximity query in a partitioned architecture
- Provide optimizations for complex algorithms, such as the DBSCAN-MR density-based clustering algorithm


Spatial-aware optimizations
Partitioning and advanced query optimizers
Geopark (including spatial representational support)
Spark Core



# Layered Spatial-aware In-memory Processing Optimization Architecture

---

- Our support **does not require to modify GeoSpark** and is a transparent layer atop of it that hides implementation details from application layer
- Our support **includes spatial-aware partitioning strategies** better trading off three challenges, such as load balancing, containment query optimization, and spatial co-locality evaluation
- Our support **includes a DBSCAN-MR implementation** (which belongs to the co-location data mining family) able to work atop GeoSpark





# Self-Adaptation

---

- ***Density-aware spatial data partitioning***, based on spatial object's distribution density
  - Useful with heavily skewed datasets
  - Roughly balances loads across computing resources
- For the execution of new clustering sessions, a ***self-tuning module optimizes cutting factors for subsequent sessions***, gaining performance improvement


# Data Partitioning




- Considering the earth flattened out, we can define ***DBSCAN-MR cutting factors as vertical partitioning lines*** in planar geometry
- ***Cutting factors configuration has an impact*** on partition's loads and spatial boundary objects




# Self-Adaptable Spatial-Aware Partitioner (SASAP)






# Self-adaptation of vertical cutting factors




- Monitors execution times at each data partition
- Uses a shared Accumulator variable
- CalculateSplits defines the new configuration, using ***threshold-based solutions to decide when to stop configuration tuning***



# Experimental Setup

- Our experimental setup utilized ***Amazon AWS cloud's computing services***, specifically ***AWS EC2*** service
- 5 nodes have been used for deployment, one master and four processing
- On each node, spark 1.6.2 was installed, and ganglia 3.7.2 was used for performance analysis. our input database consisted of 250,000 spatial objects collected through ParticipAct project
- Comparison with a MongoDB-based Map-Reduce implementation (w/out using in memory RDDs)



# Experimental Setup (cont'ed)


5 instances m3.xlarge using 1vCPU e 15GB RAM


1 Master:

Driver Spark,  
ResourceManager  
YARN and  
NameNode HDFS



4 Slaves:

Executor Spark  
and DataNode  
HDFS






# Self-Adaptable Load Balancing



First running of DBSCAN-MR ( $\approx$  50 mins) Tenth running of DBSCAN-MR ( $\approx$  20 mins)


- CPU load during running sessions 1 and 10
- For 5 nodes configuration, and a dataset consisting of 250000 spatial entries

# Query Performance Optimization




Execution time ***continuously improves*** along different processing rounds and goes **from 50.1 minutes** (first iteration) **to 22.7 minutes** (tenth iteration), obtaining a percentage of speed-up improvement equivalent to 54.7%

# GeoSpark vs MongoDB




- Performance applying our GeoSpark support compared with DBSCAN-MR implementation over MongoDB
- Speedup gain degrades as we increase the data size,



# Conclusions & Ongoing Works

---

- **Conclusions**
  - Current ***support of big data tools for compute-intensive geospatial big data*** sets is still rather ***poor***
  - Our framework efficiently supports **querying and analyzing big geospatial data** and was plugged on top of GeoSpark, with a motivation to optimize the performance of DBSCAN-MR clustering
- **Future works**
  - Incorporating additional services such as integrated geospatial-aware machine learning and data mining service
  - Extending our framework so to enable ***online processing of geospatial data streams*** in **STARK**




---

**Thank you!**

# **In-memory Spatial-Aware Framework for Processing Proximity-Alike Queries in Big Spatial Data**

**Luca Foschini**, Assistant Professor  
DISI - University of Bologna  
Viale del Risorgimento, 2 - 40136 Bologna - ITALY  
Email: [luca.foschin@unibo.it](mailto:luca.foschin@unibo.it)  
Web: <https://www.unibo.it/sitoweb/luca.foschini>



# Contacts

---

<http://lia.deis.unibo.it/Staff/LucaFoschini/>  
<http://middleware.unibo.it/>



ALMA MATER STUDIORUM  
UNIVERSITÀ DI BOLOGNA  
DIPARTIMENTO DI INFORMATICA - SCIENZA E INGEGNERIA